

## A short introduction to GPU programming



Kees Lemmens,  
Email: [C.W.J.Lemmens@tudelft.nl](mailto:C.W.J.Lemmens@tudelft.nl),  
Delft Institute for Applied Mathematics (DIAM),  
Faculty of EEMCS, Delft University of Technology,  
The Netherlands.

## Part 1: Introduction to GPUs

# Historical overview of parallel programming

1985-95: Super (vector) computers era (Convex, Cray, TM)

1995- : Beowulf concept: Linux based low budget PC's

2000- : Linux clusters: Linux based (very) high budget PC's

2004- : Vector computing on GPU cards designed for gaming

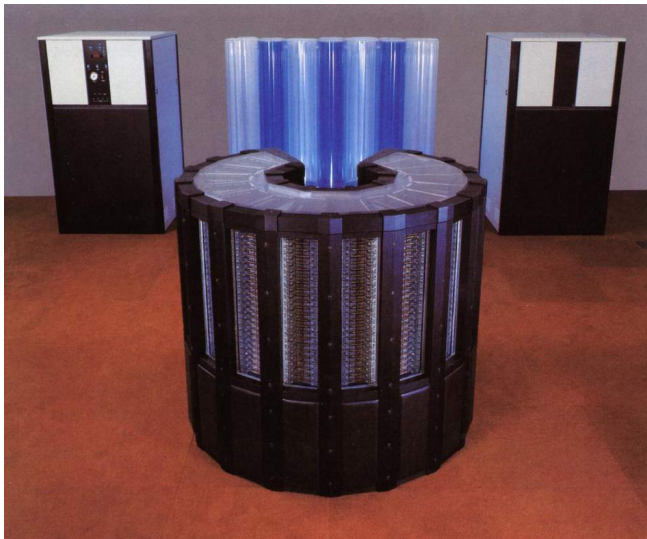
2008- : High end vector computing on dedicated GPU cards (Nvidia)

2010- : Multi GPU systems:

- Multicore systems with 4-8 GPUs
- Large GPU clusters with thousands of GPUs

2015- : Focus moving from Numerical Methods to Machine Learning and AI

## USA: NASA Cray-2 vector Computer, 1.9 Gigafllops, 250 kW (1985)



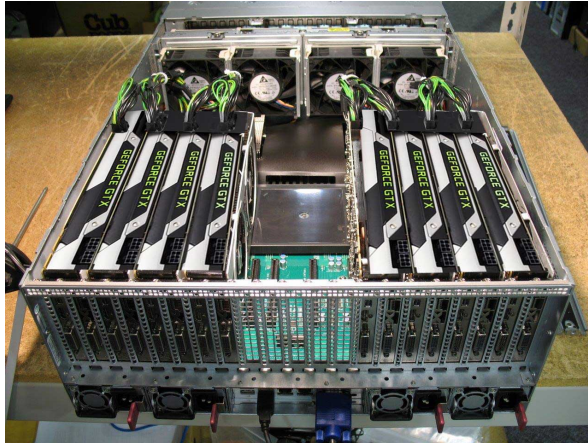
China: Tianhe-2, 34 Petaflops, 24 MWatt, 390 M\$, 38.000 Xeon + 44.000 Xeon Phi (2016)



USA: Summit, 150 Petaflops, 13 MWatt, 325 M\$, 9200 Power9 + 28.000 Tesla V100 GPU (2019)



USA: Maxwell 8 GPU desktop 20-30 Teraflops (SP),  
1.6 kWatt, 15 k\$, 8 GPUs GTX980 (2015)

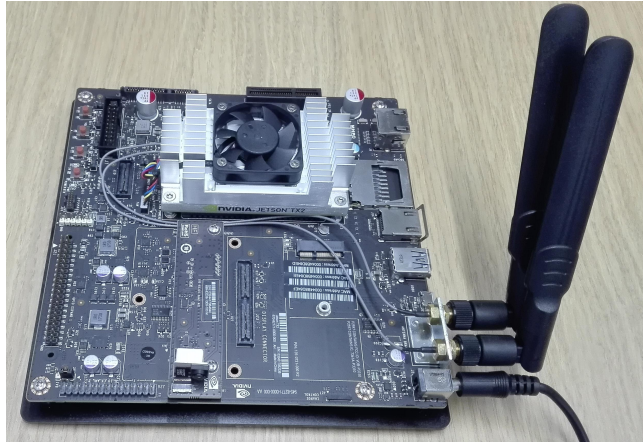


USA: Bizon 4 GPU desktop, 65 Teraflops (SP),  
1 kWatt, 3 k\$, 4 GPUs RTX8000 (2022)





USA San Jose: Jetson Tegra X2 development kit,  
30 Watt, 300 \$, 1 GPU Pascal on Tegra with 4 core ARM 64 (2017)



# Parallel computing on Graphical Processing Units 1

**GPU:** Use graphics (video) cards for scientific computing

## Features:

- Based upon standard video cards by Nvidia (or ATI)
- Uses standard computer with either Linux, MSW or MacOS
- Programming model SIMD (Single Instruction, Multiple Data)
- Parallelisation inside card through “threads”
- Dedicated software to access card and start “kernels” (Cuda, OpenCL)

Most popular software toolkits are **CUDA** by [Nvidia](#) and **OpenCL** by [Khronos](#), but we'll mainly talk about **CUDA**, as it is the most commonly used toolkit today.

## Parallel computing on Graphics Cards 2

### Advantages

- Hardware is cheap compared with Linux clusters (at least if gaming cards are used)
- Lower power consumption per Flop than CPU
- Simple GPU already available in many systems without any extra investment
- Capable of thousands to millions of parallel threads on a single GPU card
- Very fast for algorithms that can be efficiently parallelised (not that many ;-)
- Often better scalability than MPI for large problems due to faster communication
- New libraries hiding complexity to the user: **BLAS, FFTW, SPARSE, SOLVER**
- Easy integration with 3D graphics if code runs on the same card (or by using Prime)

## Parallel computing on Graphics Cards 3

### Disadvantages

- Limited amount of memory available (between 4-48 GByte)
- Memory transfers between host and GPU cause a performance bottleneck
- Often no ECC: error correcting memory (except expensive highend GPUs)
- Fast double precision GPUs are still quite expensive (P100, V100, A100)
- Slow for algorithms without enough data parallelism (e.g. scalar vector product)
- Debugging code on GPU can be complicated
- Achieving theoretically possible speedup is almost impossible
- Using Multi GPUs in a cluster is quite complex (often done with MPI or OpenMP)

## What is Cuda?

**Cuda:** means “Compute Unified Device Architecture” and is a software toolkit by Nvidia to ease the use of (Nvidia) graphics cards for scientific programming. It needs on a special video driver as well.

# What is Cuda?

**Cuda:** means “Compute Unified Device Architecture” and is a software toolkit by Nvidia to ease the use of (Nvidia) graphics cards for scientific programming. It needs on a special video driver as well.

## Features:

- Special C compiler to build code both for CPU and GPU (nvcc)
- C Language extensions (codewords) :
  - distinguish CPU and GPU functions
  - access different types of memory on the GPU
  - specify how code should be parallelized on the GPU
- Library routines for memory transfer between host and GPU
- Extra BLAS, Sparse and FFT libraries for easy porting existing code
- Dedicated to Nvidia hardware
- Mainly standard C on the GPU (limited support for C++ and Fortran)

# What is OpenCL?

**OpenCL:** means “Open Computing Language” and is a C99 like language for parallel kernels in combination with an API to get these kernels loaded and started on different types of parallel hardware (e.g. AMD, Intel, Nvidia).

# What is OpenCL?

**OpenCL:** means “Open Computing Language” and is a C99 like language for parallel kernels in combination with an API to get these kernels loaded and started on different types of parallel hardware (e.g. AMD, Intel, Nvidia).

## Features:

- Only contains a dedicated compiler for building GPU code
- So, no support for mixed (GPU + CPU) source code as in Cuda
- This makes starting kernels is more complicated than Cuda
- Data transfer between host and GPU is done using library functions
- Runs on several types of GPU (AMD, Nvidia) and is easily portable
- However, performance is **not necessarily portable** across platforms
- OpenCL is an Open Standard, while Cuda causes a vendor-lock to Nvidia
- Less mature than Cuda, often no or limited support for special GPU features



## When to use GPUs? Matrix Product comparison

Let us look at a computation intensive Matrix product  $A \times B$  in single precision.  
Note that the time is in seconds.

Type	512	1024	2048	4096	8192	16384
Blas/Atlas Xeon 3.6	0.013	0.10	0.78	6.11	48.67	388
CuBlas GPU K620	0.001	0.005	0.0348	0.250	1.843	-
CuBlas GPU GTX480	0.0016	0.007	0.036	0.222	1.508	-
CuBlas GPU Kepler K20	0.0008	0.0034	0.0155	0.085	0.542	3.94
CuBlas GPU TitanX	0.0007	0.0028	0.0118	0.058	0.327	1.96
CuBlas GPU P100	0.0004	0.0014	0.0064	0.032	0.192	1.14
Ratio Xeon/P100	32x	71x	122x	190x	253x	340x

## When to use GPUs? Matrix Product comparison

Let us look at a computation intensive Matrix product  $A \times B$  in single precision.  
Note that the time is in seconds.

Type	512	1024	2048	4096	8192	16384
Blas/Atlas Xeon 3.6	0.013	0.10	0.78	6.11	48.67	388
CuBlas GPU K620	0.001	0.005	0.0348	0.250	1.843	-
CuBlas GPU GTX480	0.0016	0.007	0.036	0.222	1.508	-
CuBlas GPU Kepler K20	0.0008	0.0034	0.0155	0.085	0.542	3.94
CuBlas GPU TitanX	0.0007	0.0028	0.0118	0.058	0.327	1.96
CuBlas GPU P100	0.0004	0.0014	0.0064	0.032	0.192	1.14
Ratio Xeon/P100	32x	71x	122x	190x	253x	340x

### Conclusions:

- ▶ GPUs are much faster than CPUs, especially for larger dimensions.
- ▶ GPUs are very limited by their main memory compared to CPUs.
- ▶ Gaming cards (GTX480) good price/performance ratio, less memory.

Note : GTX480 and K20 cards in identical hostsystems perform the same (for single prec.)

## When to use GPUs? Matrix Vector Product comparison

Now look at a less computation intensive Matrix-Vector product  $A \times b$  in single precision. Note that the time is in **milliseconds**.

Type (nc=no communication)	512	1024	2048	4096	16384	32768
Blas/Atlas Xeon 3.6	0.1	0.4	1.6	7	100	430
CuBlas GPU K620	0.3	0.9	4	14	213	-
CuBlas GPU GTX480	0.4	2	6	30	400	-
CuBlas GPU K20	0.7	1	3	10	200	800
CuBlas GPU TitanX	0.3	0.9	3	10	200	700
CuBlas GPU P100	0.17	0.44	1.5	5.6	89	350
CuBlas GPU P100 (nc)	0.05	0.05	0.09	0.17	2	7.9
Ratio Xeon/P100	0.6x	0.9x	1.1x	1.25x	1.12x	1.23x
Ratio Xeon/P100 (nc)	2x	8x	18x	41x	50x	54x

## When to use GPUs? Matrix Vector Product comparison

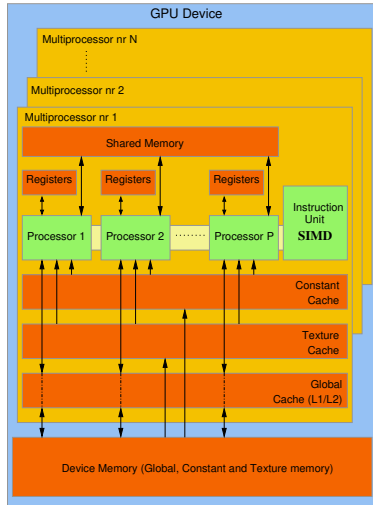
Now look at a less computation intensive Matrix-Vector product  $A \times b$  in single precision. Note that the time is in **milliseconds**.

Type (nc=no communication)	512	1024	2048	4096	16384	32768
Blas/Atlas Xeon 3.6	0.1	0.4	1.6	7	100	430
CuBlas GPU K620	0.3	0.9	4	14	213	-
CuBlas GPU GTX480	0.4	2	6	30	400	-
CuBlas GPU K20	0.7	1	3	10	200	800
CuBlas GPU TitanX	0.3	0.9	3	10	200	700
CuBlas GPU P100	0.17	0.44	1.5	5.6	89	350
CuBlas GPU P100 (nc)	0.05	0.05	0.09	0.17	2	7.9
Ratio Xeon/P100	0.6x	0.9x	1.1x	1.25x	1.12x	1.23x
Ratio Xeon/P100 (nc)	2x	8x	18x	41x	50x	54x

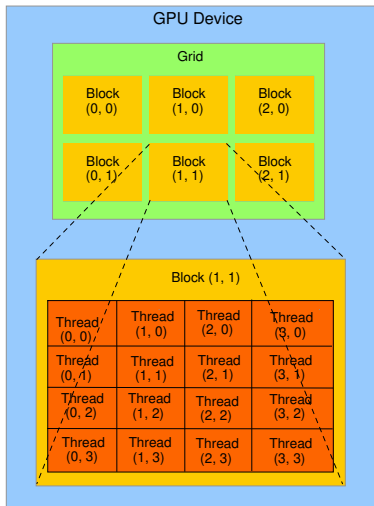
### Conclusions :

- ▶ GPU performance doesn't improve much for larger dimensions.
- ▶ Difference in performance for GPUs is only small.
- ▶ Blas on CPU is faster than on GPU unless no host transfers are required (nc).

# Hardware layout of the GPU



# Processing layout of a GPU kernel



## Comparison to a Pasta Machine (9 threads per block ;-)



Relation software  $\longleftrightarrow$  hardware

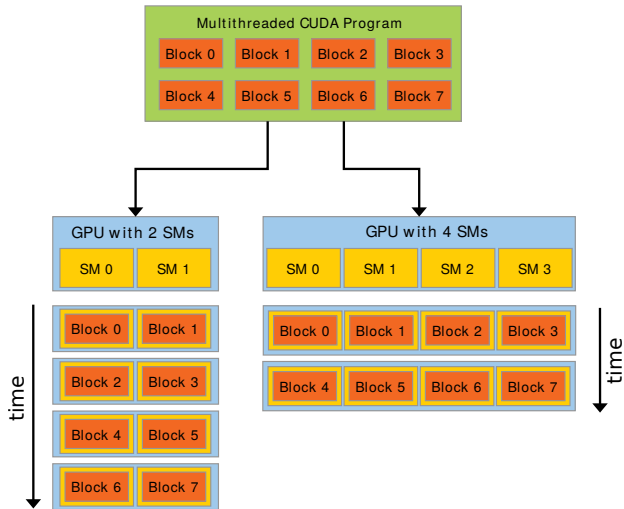
Thread  $\longleftrightarrow$  Core

Block  $\longleftrightarrow$  Multiprocessor

Grid  $\longleftrightarrow$  Device



# Scalability of a GPU kernel



## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)

## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)
- ▶ A single device can have several **streamprocessors** (SM) (K620 = 3, K20 = 13, K80 = 2x13, Titan X = 24, P100 = 28, V100=40)

## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)
- ▶ A single device can have several **streamprocessors** (SM) (K620 = 3, K20 = 13, K80 = 2x13, Titan X = 24, P100 = 28, V100=40)
- ▶ Each SM has 32 (2.0), 48 (2.x), 192 (3.x) or 128 (5-8.x) cores. A SM can run 32 active (semi) concurrent **threads** called **warps**

## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)
- ▶ A single device can have several **streamprocessors** (SM) (K620 = 3, K20 = 13, K80 = 2x13, Titan X = 24, P100 = 28, V100=40)
- ▶ Each SM has 32 (2.0), 48 (2.x), 192 (3.x) or 128 (5-8.x) cores. A SM can run 32 active (semi) concurrent **threads** called **warps**
- ▶ Threads in a SM run concurrently (max. 8-32) or sequentially if more. Maximum resident threads per SM is 1536 (2.x and up).

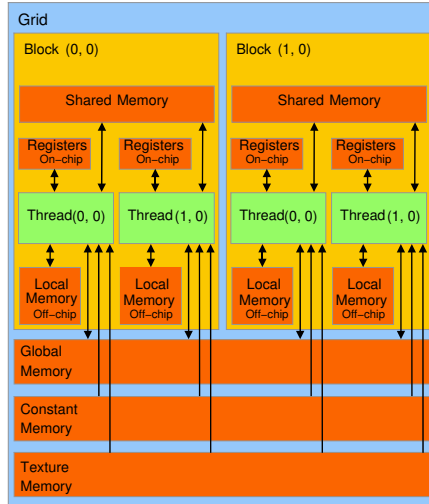
## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)
- ▶ A single device can have several **streamprocessors** (SM) (K620 = 3, K20 = 13, K80 = 2x13, Titan X = 24, P100 = 28, V100=40)
- ▶ Each SM has 32 (2.0), 48 (2.x), 192 (3.x) or 128 (5-8.x) cores. A SM can run 32 active (semi) concurrent **threads** called **warps**
- ▶ Threads in a SM run concurrently (max. 8-32) or sequentially if more. Maximum resident threads per SM is 1536 (2.x and up).
- ▶ A **block** is a set of threads running on a single SM. Maximum number of threads per block is 512-1024 (x and y) and 64 (z). The total maximum ( $x*y*z$ ) is 1024 (2.x and up)

## Some capabilities of an Nvidia GPU

- ▶ An Nvidia GPU card can have 1 or more **devices** (Tesla S1070 = 4, Tesla C2070 = 1, NVS290 = 1, K20 = 1, K80 = 4, Titan X = 1)
- ▶ A single device can have several **streamprocessors** (SM) (K620 = 3, K20 = 13, K80 = 2x13, Titan X = 24, P100 = 28, V100=40)
- ▶ Each SM has 32 (2.0), 48 (2.x), 192 (3.x) or 128 (5-8.x) cores. A SM can run 32 active (semi) concurrent **threads** called **warps**
- ▶ Threads in a SM run concurrently (max. 8-32) or sequentially if more. Maximum resident threads per SM is 1536 (2.x and up).
- ▶ A **block** is a set of threads running on a single SM. Maximum number of threads per block is 512-1024 (x and y) and 64 (z). The total maximum ( $x*y*z$ ) is 1024 (2.x and up)
- ▶ Maximum number of blocks is  $2^{31} - 1$  for each dimension (for arch 2.x  $2^{16} - 1$ ) independent of SM count. Arch 2.x and up have 3 dimensions (so x, y and z)

# Memory organisation of a GPU 1





## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory

## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory
- ▶ Global memory accessible by all streaming multiprocessors and is persistent

## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory
- ▶ Global memory accessible by all streaming multiprocessors and is persistent
- ▶ Constant and texture memory can only be *read* and have a **local cache** on each SM

## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory
- ▶ Global memory accessible by all streaming multiprocessors and is persistent
- ▶ Constant and texture memory can only be *read* and have a **local cache** on each SM
- ▶ Each SM has **shared memory** (on-chip) that can be used by all threads in an active block but is non-persistent (kind of cache)

## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory
- ▶ Global memory accessible by all streaming multiprocessors and is persistent
- ▶ Constant and texture memory can only be *read* and have a **local cache** on each SM
- ▶ Each SM has **shared memory** (on-chip) that can be used by all threads in an active block but is non-persistent (kind of cache)
- ▶ Each SM has **local memory** (off-chip) and **registers** (on-chip) that can only be used by a single thread and which is also non-persistent

## Memory organisation of a GPU 2

- ▶ Device memory consists of **global**, **constant** and **texture** memory
- ▶ Global memory accessible by all streaming multiprocessors and is persistent
- ▶ Constant and texture memory can only be *read* and have a **local cache** on each SM
- ▶ Each SM has **shared memory** (on-chip) that can be used by all threads in an active block but is non-persistent (kind of cache)
- ▶ Each SM has **local memory** (off-chip) and **registers** (on-chip) that can only be used by a single thread and which is also non-persistent
- ▶ Note that memory access times also depend on the access method! **Coalesced** access may give a large performance boost, especially for less computation intensive problems
- ▶ Coalesced access is likely much better if the number of threads per block is always a multiple of the warp size (32 threads)

## New features for GPUs architectures

- ▶ **1.x or Tesla:** Now obsolete, name still used as a brandname.
- ▶ **2.x or Fermi:** 48 cores per SM, atomic operations. L1 cache (combined with Shared Memory) on each SM and a single L2 cache per device for global memory. Explicit use of shared memory less important.

## New features for GPUs architectures

- ▶ **1.x or Tesla:** Now obsolete, name still used as a brandname.
- ▶ **2.x or Fermi:** 48 cores per SM, atomic operations. L1 cache (combined with Shared Memory) on each SM and a single L2 cache per device for global memory. Explicit use of shared memory less important.
- ▶ **3.x or Kepler:** 192 cores per SM, L1 cache only for local memory, Only L2 cache for global memory. Unified Memory Access (no explicit host- device transfers). Explicit use of shared memory again more important.



## New features for GPUs architectures

- ▶ **1.x or Tesla:** Now obsolete, name still used as a brandname.
- ▶ **2.x or Fermi:** 48 cores per SM, atomic operations. L1 cache (combined with Shared Memory) on each SM and a single L2 cache per device for global memory. Explicit use of shared memory less important.
- ▶ **3.x or Kepler:** 192 cores per SM, L1 cache only for local memory, Only L2 cache for global memory. Unified Memory Access (no explicit host- device transfers). Explicit use of shared memory again more important.
- ▶ **5.x or Maxwell:** 128 cores per SM, L1 cache independent from Shared Memory. Dynamic Parallellism: start kernels from a running kernel.

## New features for GPUs architectures

- ▶ **1.x or Tesla:** Now obsolete, name still used as a brandname.
- ▶ **2.x or Fermi:** 48 cores per SM, atomic operations. L1 cache (combined with Shared Memory) on each SM and a single L2 cache per device for global memory. Explicit use of shared memory less important.
- ▶ **3.x or Kepler:** 192 cores per SM, L1 cache only for local memory, Only L2 cache for global memory. Unified Memory Access (no explicit host- device transfers). Explicit use of shared memory again more important.
- ▶ **5.x or Maxwell:** 128 cores per SM, L1 cache independent from Shared Memory. Dynamic Parallellism: start kernels from a running kernel.
- ▶ **6.x or Pascal:** 128 cores per SM, Half precision FP (16-bit), NVLink bus to directly link GPUs, bypassing PCI-E bus (5-15 times faster).

## New features for GPUs architectures

- ▶ **1.x or Tesla**: Now obsolete, name still used as a brandname.
- ▶ **2.x or Fermi**: 48 cores per SM, atomic operations. L1 cache (combined with Shared Memory) on each SM and a single L2 cache per device for global memory. Explicit use of shared memory less important.
- ▶ **3.x or Kepler**: 192 cores per SM, L1 cache only for local memory, Only L2 cache for global memory. Unified Memory Access (no explicit host- device transfers). Explicit use of shared memory again more important.
- ▶ **5.x or Maxwell**: 128 cores per SM, L1 cache independent from Shared Memory. Dynamic Parallellism: start kernels from a running kernel.
- ▶ **6.x or Pascal**: 128 cores per SM, Half precision FP (16-bit), NVLink bus to directly link GPUs, bypassing PCI-E bus (5-15 times faster).
- ▶ **7.0 or Volta**: more half precision FP, Tensor cores (deep learning).
- ▶ **7.5 or Turing**: mainly graphic and raytracing improvements.
- ▶ **8.0 or Ampere**: several Tensor core improvements.

## Example of parallel work: Amish building a barn



## Example of parallel work: Amish raising a barn 2



## Example of parallel work

If a single Amish can build a barn in 50 days then  
50 Amish can build that barn in a single day!

## Example of parallel work

If a single Amish can build a barn in 50 days then  
50 Amish can build that barn in a single day!

However, in case of ordinary people **and** for GPUs  
this is far to optimistic ...

## Parallel matrix summation analysis

Suppose we need the **sum** of 2 (large) matrices  $A$  and  $B$  :

$$\begin{pmatrix} A(1,1) & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & A(N,N) \end{pmatrix} + \begin{pmatrix} B(1,1) & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & B(N,N) \end{pmatrix} = C$$

How could we parallelize this and what would be the speedup?



## Parallel matrix summation: speedup using strips

Sequential calculation:  $N.N$  elements, 1 operation/element:

$$T_{seq} = N^2.t_1 = O(N^2)$$

## Parallel matrix summation: speedup using strips

Sequential calculation:  $N.N$  elements, 1 operation/element:

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N \cdot \frac{N}{P}$  elements, 1 op/el:

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

## Parallel matrix summation: speedup using strips

Sequential calculation:  $N.N$  elements, 1 operation/element:

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N \cdot \frac{N}{P}$  elements, 1 op/el:

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A + B$  + gather  $C$  :

$$T_{com} = P \cdot 3 \cdot \frac{N^2}{P} \cdot t_2 = O(N^2)$$

## Parallel matrix summation: speedup using strips

Sequential calculation:  $N.N$  elements, 1 operation/element:

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N \cdot \frac{N}{P}$  elements, 1 op/el:

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A + B$  + gather  $C$  :

$$T_{com} = P \cdot 3 \cdot \frac{N^2}{P} \cdot t_2 = O(N^2)$$

$$\text{For } (P > \frac{t_1}{t_2}) : T_{par} = T_{cal} + T_{com} = O\left(\frac{N^2}{P}\right) + O(N^2) \approx O(N^2)$$

## Parallel matrix summation: speedup using strips

Sequential calculation:  $N.N$  elements, 1 operation/element:

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N \cdot \frac{N}{P}$  elements, 1 op/el:

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A + B$  + gather  $C$  :

$$T_{com} = P \cdot 3 \cdot \frac{N^2}{P} \cdot t_2 = O(N^2)$$

$$\text{For } (P > \frac{t_1}{t_2}) : T_{par} = T_{cal} + T_{com} = O\left(\frac{N^2}{P}\right) + O(N^2) \approx O(N^2)$$

$$\text{Speedup : } \frac{T_{seq}}{T_{par}} = \frac{O(N^2)}{O(N^2)} = O(1), \text{ some for } P < \frac{t_1}{3 \cdot t_2} \text{ (no } N \text{ dep. !)}$$

## Parallel matrix vector product analysis

Suppose we need the **product** of a (large) matrix  $A$  and vector  $x$  :

$$\begin{pmatrix} A(1,1) & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & A(N,N) \end{pmatrix} * \begin{pmatrix} x(1) \\ \dots \\ x(N) \end{pmatrix} = b$$

How could we parallelize this and what would be the speedup?

## Parallel matrix vector product analysis: speedup using strips

Sequential calculation:  $N$  computations for each element of  $x$  of size  $N$  :

$$T_{seq} = N^2.t_1 = O(N^2)$$

## Parallel matrix vector product analysis: speedup using strips

Sequential calculation:  $N$  computations for each element of  $x$  of size  $N$  :

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N$  computations for  $\frac{N}{P}$  elements of  $x$  :

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$



## Parallel matrix vector product analysis: speedup using strips

Sequential calculation:  $N$  computations for each element of  $x$  of size  $N$  :

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N$  computations for  $\frac{N}{P}$  elements of  $x$  :

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A$  +  $x$  + gather  $b$  :

$$T_{com} = P \cdot \left( \frac{N^2}{P} + 2 \cdot \frac{N}{P} \right) \cdot t_2 = O(N^2)$$

## Parallel matrix vector product analysis: speedup using strips

Sequential calculation:  $N$  computations for each element of  $x$  of size  $N$  :

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N$  computations for  $\frac{N}{P}$  elements of  $x$  :

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A$  +  $x$  + gather  $b$  :

$$T_{com} = P \cdot \left( \frac{N^2}{P} + 2 \cdot \frac{N}{P} \right) \cdot t_2 = O(N^2)$$

$$\text{For } (P \gg \frac{t_1}{t_2}) : T_{par} = T_{cal} + T_{com} = O\left(\frac{N^2}{P}\right) + O(N^2) \approx O(N^2)$$

## Parallel matrix vector product analysis: speedup using strips

Sequential calculation:  $N$  computations for each element of  $x$  of size  $N$  :

$$T_{seq} = N^2 \cdot t_1 = O(N^2)$$

Parallel calculation for strips:  $N$  computations for  $\frac{N}{P}$  elements of  $x$  :

$$T_{cal} = \frac{N^2}{P} \cdot t_1 = O\left(\frac{N^2}{P}\right)$$

Parallel communication for strips: scatter  $A$  +  $x$  + gather  $b$  :

$$T_{com} = P \cdot \left( \frac{N^2}{P} + 2 \cdot \frac{N}{P} \right) \cdot t_2 = O(N^2)$$

$$\text{For } (P \gg \frac{t_1}{t_2}) : T_{par} = T_{cal} + T_{com} = O\left(\frac{N^2}{P}\right) + O(N^2) \approx O(N^2)$$

$$\text{Speedup : } \frac{T_{seq}}{T_{par}} = \frac{O(N^2)}{O(N^2)} = O(1), \text{ some speedup for } P < \frac{t_1}{t_2} \text{ (no } N \text{ dep. !)}$$

## Parallel matrix multiplication analysis

Suppose we need the **product** of 2 (large) matrices  $A$  and  $B$ :

$$\begin{pmatrix} A(1,1) & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & A(N,N) \end{pmatrix} * \begin{pmatrix} B(1,1) & \dots & \dots \\ \dots & \dots & \dots \\ \dots & \dots & B(N,N) \end{pmatrix} = C$$

How could we parallelize this and what would be the speedup?

## Parallel matrix multiplication: speedup using strips

Sequential calculation :  $N.N$  elements,  $N$  operations/element:

$$T_{seq} = N^3.t_1 = O(N^3)$$

## Parallel matrix multiplication: speedup using strips

Sequential calculation :  $N.N$  elements,  $N$  operations/element:

$$T_{seq} = N^3 \cdot t_1 = O(N^3)$$

Parallel calculation for strips :  $N \cdot \frac{N}{P}$  elements ,  $N$  op/el:

$$T_{cal} = \frac{N^3}{P} \cdot t_1 = O\left(\frac{N^3}{P}\right)$$

## Parallel matrix multiplication: speedup using strips

Sequential calculation :  $N.N$  elements,  $N$  operations/element:

$$T_{seq} = N^3 \cdot t_1 = O(N^3)$$

Parallel calculation for strips :  $N \cdot \frac{N}{P}$  elements ,  $N$  op/el:

$$T_{cal} = \frac{N^3}{P} \cdot t_1 = O\left(\frac{N^3}{P}\right)$$

Parallel communication for strips : scatter  $A$  + broadcast  $B$  + gather  $C$  :

$$T_{com} = P \cdot \left(2 \cdot \frac{N^2}{P} + N^2\right) \cdot t_2 \approx O(P \cdot N^2)$$

## Parallel matrix multiplication: speedup using strips

Sequential calculation :  $N.N$  elements,  $N$  operations/element:

$$T_{seq} = N^3 \cdot t_1 = O(N^3)$$

Parallel calculation for strips :  $N \cdot \frac{N}{P}$  elements ,  $N$  op/el:

$$T_{cal} = \frac{N^3}{P} \cdot t_1 = O\left(\frac{N^3}{P}\right)$$

Parallel communication for strips : scatter  $A$  + broadcast  $B$  + gather  $C$  :

$$T_{com} = P \cdot \left(2 \cdot \frac{N^2}{P} + N^2\right) \cdot t_2 \approx O(P \cdot N^2)$$

$$\text{For } (N \gg P \frac{t_2}{t_1}): T_{par} = T_{cal} + T_{com} = O\left(\frac{N^3}{P}\right) + O(P \cdot N^2) \approx O\left(\frac{N^3}{P}\right)$$



## Parallel matrix multiplication: speedup using strips

Sequential calculation :  $N.N$  elements,  $N$  operations/element:

$$T_{seq} = N^3 \cdot t_1 = O(N^3)$$

Parallel calculation for strips :  $N \cdot \frac{N}{P}$  elements ,  $N$  op/el:

$$T_{cal} = \frac{N^3}{P} \cdot t_1 = O\left(\frac{N^3}{P}\right)$$

Parallel communication for strips : scatter  $A$  + broadcast  $B$  + gather  $C$  :

$$T_{com} = P \cdot \left(2 \cdot \frac{N^2}{P} + N^2\right) \cdot t_2 \approx O(P \cdot N^2)$$

$$\text{For } (N \gg P \frac{t_2}{t_1}): T_{par} = T_{cal} + T_{com} = O\left(\frac{N^3}{P}\right) + O(P \cdot N^2) \approx O\left(\frac{N^3}{P}\right)$$

$$\text{Speedup : } \frac{T_{seq}}{T_{par}} = \frac{O(N^3)}{O\left(\frac{N^3}{P}\right)} = O(P) , \text{ speedup as long as: } P^2 < \frac{N \cdot t_1}{t_2}$$

## Part 2: Hands on GPU, practical aspects

## Control of GPU under CUDA and C: compilers and linkers

To make a Cuda program for the GPU we need some special statements in the code and extra options during compiling and linking.

CUDA provides its own compiler for C: **nvcc**. It compiles both normal CPU code (for which it simply uses the native C-compiler) and GPU code.

## Control of GPU under CUDA and C: compilers and linkers

To make a Cuda program for the GPU we need some special statements in the code and extra options during compiling and linking.

CUDA provides its own compiler for C: **nvcc**. It compiles both normal CPU code (for which it simply uses the native C-compiler) and GPU code.

Some things to remember:

- ▶ Manually compile using : `% nvcc prg.c -o prg`
- ▶ Automatically compile by using e.g. a `Makefile`
- ▶ GPU code and CPU code can be in the same source (extension `.cu`)

# Parallel “Hello World1” in Cuda C

A simple **Hello World** example in C.

```
#include <stdio.h>
#include <cuda.h>

#define NRBLKS    4 // Nr of blocks in a kernel (gridDim)
#define NRTPBK    4 // Nr of threads in a block (blockDim)

__global__ void printOnGPU()
{
    int myid = (blockIdx.x * blockDim.x) + threadIdx.x;
    printf("Hello World, I am thread %d \n",myid);
}

int main(void)
{
    printOnGPU <<< NRBLKS, NRTPBK >>> ();
    cudaDeviceReset(); // Without this line the output is NOT shown !
}
```

# Parallel “Hello World2” in Cuda C

```
#include <stdio.h>
#include <cuda.h>
#define NRBLKS    4 // Nr of blocks in a kernel (gridDim)
#define NRTPBK    4 // Nr of threads in a block (blockDim)

__global__ void decodeOnGPU(char *string)
{ int myid = (blockIdx.x * blockDim.x) + threadIdx.x;
  string[myid] ^= (char)011; // 000001001 = 011 in octal
}

int main(void)
{ char *encryption = "Aleef"^(f{em})(\11";
  char *string_h, * string_d; // pointers to host and device memory
  int len = strlen(encryption);

  string_h = (char *)malloc(    len);
  cudaMalloc((void **)&string_d, len);
  cudaMemcpy(string_d, encryption, len, cudaMemcpyHostToDevice);
  decodeOnGPU <<< NRBLKS, NRTPBK >>> (string_d);
  cudaMemcpy(string_h, string_d, len, cudaMemcpyDeviceToHost);

  printf("%s\n",string_h);
  cudaFree(string_d);
  free(string_h);
}
```

# Unified Memory

A new feature since architecture 3.0 is **Unified Memory**. This creates a **single virtual memory space** for all GPU(s) and host(s) memory and automatically copies data if accessed from a place where it is not physically located. As such it makes `cudaMemcpy()` calls unnecessary, although the actual memory copies still need to be done under the hood.

## Points of attention

- ▶ Keywords: `CudaMallocManaged()` or `__device__ __managed__`
- ▶ Can be slower than explicit memory management (useless transfers)...
- ▶ Must be compiled with architecture `sm_30` or higher
- ▶ Some Intel hardware MMUs may confuse this function, causing wrong answers without warnings. Boot with `iommu=soft` (on Linux)

# Unified Memory

A new feature since architecture 3.0 is **Unified Memory**. This creates a **single virtual memory space** for all GPU(s) and host(s) memory and automatically copies data if accessed from a place where it is not physically located. As such it makes `cudaMemcpy()` calls unnecessary, although the actual memory copies still need to be done under the hood.

## Points of attention

- ▶ Keywords: `CudaMallocManaged()` or `__device__ __managed__`
- ▶ Can be slower than explicit memory management (useless transfers)...
- ▶ Must be compiled with architecture `sm_30` or higher
- ▶ Some Intel hardware MMUs may confuse this function, causing wrong answers without warnings. Boot with `iommu=soft` (on Linux)

**Example:** Discuss and run `hello_unifiedmemory`



# Parallel “Hello World2” in Cuda C (Unified Memory)

```
#include <stdio.h>
#include <cuda.h>

#define NRBLKS    4 // Nr of blocks in a kernel (gridDim)
#define NRTPBK    4 // Nr of threads in a block (blockDim)

__device__ __managed__ char string[] = "Aleef)^f{em)((\11";

__global__ void decodeOnGPU(char *string)
{ int myid = (blockIdx.x * blockDim.x) + threadIdx.x;
  string[myid] ^= (char)011; // 00001001 = 011 in octal
}

int main(void)
{
  decodeOnGPU <<< NRBLKS, NRTPBK >>> (string);
  cudaDeviceSynchronize();

  printf("%s\n",string);
}
```

# Generic Cuda program framework

```
#include <stdio.h>
#include <cuda.h>

__global__ void routineOnGPU(float *data, int N)
{
    <GPU code>;
}

int main(void)
{
    float *data_h; // pointer to host    memory
    float *data_d; // pointer to device memory
    int size = N * sizeof(float);

    data_h = (float *)malloc(size);
    cudaMalloc((void **) &data_d, size);

    cudaMemcpy(data_d, data_h, size, cudaMemcpyHostToDevice);
    routineOnGPU <<< numberOfBlocks, threadsPerBlock >>> (data_d, N);
    cudaMemcpy(data_h, data_d, size, cudaMemcpyDeviceToHost);

    cudaFree(data_d);
    free(data_h);
}
```

# Generic Cuda program framework with Unified Memory

```
#include <stdio.h>
#include <cuda.h>

__global__ void routineOnGPU(float *data, int N)
{
    <GPU code>;
}

int main(void)
{
    float *data; // pointer to Unified host AND device memory
    int size = N * sizeof(float);

    cudaMallocManaged((void **) &data, size);

    routineOnGPU <<< numberOfBlocks, threadsPerBlock >>> (data, N);
    cudaDeviceSynchronize(); // or else no access to "data" from host

    cudaFree(data);
}
```

## Another example: Scalar-Vector product (Multiply)

**Example:** Discussion of a Scalar-Vector product with various amounts of blocks. multiply1 uses simply 1 thread per block, multiply2 uses more.

## Another example: Scalar-Vector product (Multiply)

**Example:** Discussion of a Scalar-Vector product with various amounts of blocks. `multiply1` uses simply 1 thread per block, `multiply2` uses more.

**Test:** Let's change the multiplication in `multiply3.cu` into the more time-consuming computation of  $\sin(\sqrt{a[i]} * s)$  (changing the  $t1/t2$  ratio) and recompile using `make`. Note that we use single precision to keep the results comparable: `sinf(sqrtf(a[i] * s))`

## Another example: Scalar-Vector product (Multiply)

**Example:** Discussion of a Scalar-Vector product with various amounts of blocks. `multiply1` uses simply 1 thread per block, `multiply2` uses more.

**Test:** Let's change the multiplication in `multiply3.cu` into the more time-consuming computation of  $\sin(\sqrt{a[i]} * s)$  (changing the  $t1/t2$  ratio) and recompile using `make`. Note that we use single precision to keep the results comparable: `sinf(sqrtf(a[i] * s))`

**Demo:** Let's check for both `multiply2` and 3 how the ratio between sequential and parallel changes for  $N=1000$  ( $10^3$ ),  $N=100000$  ( $10^5$ ) and  $N=10000000$  ( $10^7$ ) .

## Scalar-Vector product (Multiply): K620 vs. Xeon 3.7 GHz

Problem size	$N=10^3$	$N=10^5$	$N=10^7$
Sequential multiply2 (ms)	2	20	3400
Parallel multiply2 (ms)	60	320	16000
ratio Parallel/Sequential	30	16	5

## Scalar-Vector product (Multiply): K620 vs. Xeon 3.7 GHz

Problem size	$N=10^3$	$N=10^5$	$N=10^7$
Sequential multiply2 (ms)	2	20	3400
Parallel multiply2 (ms)	60	320	16000
ratio Parallel/Sequential	30	16	5

Problem size	$N=10^3$	$N=10^5$	$N=10^7$
Sequential multiply3 (ms)	25	1200	123000
Parallel multiply3 (ms)	74	330	16000
ratio Parallel/Sequential	3	0.27	0.13



## Starting a kernel on the GPU

Recall the line from the previous examples:

```
routineOnGPU <<<numberOfBlocks, threadsPerBlock>>>(data_d);
```

Apart from this there are 2 extra and optional arguments:

- ▶ Reserve shared memory (often less important if caching is available).
- ▶ Associated stream in case more than 1 stream is used.

## Starting a kernel on the GPU

Recall the line from the previous examples:

```
routineOnGPU <<<numberOfBlocks, threadsPerBlock>>>(data_d);
```

Apart from this there are 2 extra and optional arguments:

- ▶ Reserve shared memory (often less important if caching is available).
- ▶ Associated stream in case more than 1 stream is used.

Up to now we used 2 integers for number of blocks and threads per block. But generally these are of type `dim3` : a 3 dimensional vector with members `dim3.x`, `dim3.y` and `dim3.z`

## Starting a kernel on the GPU

Recall the line from the previous examples:

```
routineOnGPU <<<numberOfBlocks, threadsPerBlock>>>(data_d);
```

Apart from this there are 2 extra and optional arguments:

- ▶ Reserve shared memory (often less important if caching is available).
- ▶ Associated stream in case more than 1 stream is used.

Up to now we used 2 integers for number of blocks and threads per block. But generally these are of type `dim3` : a 3 dimensional vector with members `dim3.x`, `dim3.y` and `dim3.z`

**Max nr of blocks:**  $2^{31} - 1$  for x,y and z dims ( $2^{16} - 1$  for x,y on old archs).

**Max nr of threads:** 1024 for x,y and 64 for z (512 for x,y on old archs).

# CUDA: get device properties at runtime

To optimise a program for a certain GPU it is often required to obtain runtime information. This is done with `cudaGetDeviceProperties`:

```
int t,devcount;          // holds number of GPUs
cudaDeviceProp prop;     // holds GPU information

cudaGetDeviceCount(&devcount);

for(t=0;t < devcount; t++)
{ cudaGetDeviceProperties(&prop, t);
  fprintf(stderr,"memory   : %ld Bytes\n",prop.totalGlobalMem);
  fprintf(stderr,"major id: %d \n",prop.major);
  fprintf(stderr,"minor id: %d \n",prop.minor);
}
```

**Example:** Discussion of `DeviceInfo-GPU`

## CUDA high level libraries: the CUBLAS library

Cuda also contains a BLAS library that works much like the popular **BLAS** library, but a great deal faster, especially for level 3 routines. Most of the gory GPU stuff is done inside the library and is something you don't have to think about.

**Example:** Discussion of [MatrixProd-Blas-GPU](#) (framework on next slide)

## CUDA high level libraries: the CUBLAS library

Cuda also contains a BLAS library that works much like the popular **BLAS** library, but a great deal faster, especially for level 3 routines. Most of the gory GPU stuff is done inside the library and is something you don't have to think about.

**Example:** Discussion of [MatrixProd-Blas-GPU](#) (framework on next slide)

**cublasSgemm**

## CUDA high level libraries: the CUBLAS library

Cuda also contains a BLAS library that works much like the popular **BLAS** library, but a great deal faster, especially for level 3 routines. Most of the gory GPU stuff is done inside the library and is something you don't have to think about.

**Example:** Discussion of `MatrixProd-Blas-GPU` (framework on next slide)

### `cublasSgemm`

**Demo:** Show the CPU (Gsl version) and GPU programs for  $N \times N$  matrices with  $N=1024$ ,  $N=4096$ ,  $N=2048$ ,  $N=2047$  and  $N=2049$  and compare the results.

## CUDA high level libraries: the CUBLAS library

Cuda also contains a BLAS library that works much like the popular **BLAS** library, but a great deal faster, especially for level 3 routines. Most of the gory GPU stuff is done inside the library and is something you don't have to think about.

**Example:** Discussion of `MatrixProd-Blas-GPU` (framework on next slide)

### `cublasSgemm`

**Demo:** Show the CPU (Gsl version) and GPU programs for  $N \times N$  matrices with  $N=1024$ ,  $N=4096$ ,  $N=2048$ ,  $N=2047$  and  $N=2049$  and compare the results.

**Note:** The difference between 2047 and 2048 can be almost a factor 4 on older GPU architectures! However, on 2.x and up and on CPU there should hardly be any difference.



# Generic Cuda Blas framework Sgemm

```
#include <stdio.h>
#include <cuda.h>
#include <cublas_v2.h>
int main(void)
{
    float *data_hA, *data_hB, *data_hC; // pointers to host memory
    float *data_dA, *data_dB, *data_dC; // pointers to device memory
    cublasHandle_t handle;
    int size = dim * dim * sizeof(float);

    cublasCreate(&handle);
    data_hA = (float *)malloc(size); // 3 times for A, B and C
    cudaMalloc((void **) &data_dA, size); // 3 times for A, B and C

    cublasSetVector(dim*dim, sizeof(float), data_hA, 1, data_dA, 1); // 2x : A and B
    status = cublasSgemm(handle, ... data_dA, dim, data_dB, dim, ..., data_dC, dim);
    cublasGetVector(dim*dim, sizeof(float), data_dC, 1, data_hC, 1);

    cublasDestroy(handle);
    cudaFree(data_dA); // 3 times for A, B and C
    free(data_hA); // 3 times for A, B and C
}
```

# Generic Cuda Blas framework Sgemv

```
#include <cuda.h>
#include <cublas_v2.h>
int main(void)
{
    float *h_A, *h_x, *h_b; // matrix A and vectors b,x in host memory
    float *d_A, *d_x, *d_b; // matrix A and vectors b,x in device memory
    float alpha = 1, beta = 0;
    cublasHandle_t handle;

    cublasCreate(&handle);
    cudaMallocHost(&h_A, dim * dim * sizeof(float));
    cudaMallocHost(&h_x, dim * sizeof(float)); // and once more for b
    cudaMalloc(&d_A, dim * dim * sizeof(float));
    cudaMalloc(&d_x, dim * sizeof(float)); // and once more for b

    cublasSetVector(dim*dim, sizeof(float), h_A, 1, d_A, 1);
    cublasSetVector(dim, sizeof(float), h_x, 1, d_x, 1);
    status = cublasSgemv(handle, ... , &alpha, d_A, dim, d_x, 1, &beta, d_b, 1);
    cublasGetVector(dim, sizeof(float), d_b, 1, h_b, 1);

    cublasDestroy(handle);
    cudaFree(d_A); cudaFree(d_x); cudaFree(d_b);
    cudaFreeHost(h_A); cudaFreeHost(h_x); cudaFreeHost(h_b);
}
```

# Generic Cuda Blas framework Sgemv (unified memory)

```
#include <cuda.h>
#include <cublas_v2.h>
int main(void)
{
    float *h_A, *h_x, *h_b; // matrix A and vectors b,x in managed memory
    float alpha = 1, beta = 0;
    cublasHandle_t handle;

    cublasCreate(&handle);
    cudaMallocManaged(&h_A, dim * dim * sizeof(float));
    cudaMallocManaged(&h_x, dim * sizeof(float)); // and once more for b

    status = cublasSgemv(handle, ... , &alpha, h_A, dim, h_x, 1, &beta, h_b, 1);
    cudaDeviceSynchronize(); // must wait until result can be used !

    cublasDestroy(handle);
    cudaFree(h_A); cudaFree(h_x); cudaFree(h_b);
}
```

## Debugging of memory leak errors using cuda-memcheck

`cuda-memcheck` is a memory leak checker for CUDA and works much the same as the Unix tool `valgrind`. Note that it is replaced since cuda 11 with `compute-sanitizer`, which however works in exactly the same way.

### Notes

- ▶ Start it with: `% compute-sanitizer <yourprogram>`
- ▶ Compute-sanitizer always stops after the first memory error. You can force it to continue for the rest of the program using:  
`% compute-sanitizer --destroy-on-device-error kernel <yourprogram>`

# Debugging of memory leak errors using cuda-memcheck

`cuda-memcheck` is a memory leak checker for CUDA and works much the same as the Unix tool `valgrind`. Note that it is replaced since cuda 11 with `compute-sanitizer`, which however works in exactly the same way.

## Notes

- ▶ Start it with: `% compute-sanitizer <yourprogram>`
- ▶ Compute-sanitizer always stops after the first memory error. You can force it to continue for the rest of the program using:  
`% compute-sanitizer --destroy-on-device-error kernel <yourprogram>`

**Demo:** Let's check the `Memcheck-GPU` example using `compute-sanitizer` and see if you can find what is wrong in the code there?

## Profiling Cuda code

Since Cuda 5 there is a commandline profiler `nvprof`, now being replaced (Cuda 11) with `nsys nvprof`. Use as: `nvprof myprogram <arguments>`

There are also GUI profilers named `nvvp` since Cuda 4 and one named `nsys-ui` since Cuda 11 that replaces it.

**Example:** `matrixprod-simple` with 2 kernels using `nvprof` (arch  $\leq 7.x$ ).

## Profiling Cuda code

Since Cuda 5 there is a commandline profiler **nvprof**, now being replaced (Cuda 11) with **nsys nvprof**. Use as: **nvprof myprogram <arguments>**

There are also GUI profilers named **nvvp** since Cuda 4 and one named **nsys-ui** since Cuda 11 that replaces it.

**Example:** **matrixprod-simple** with 2 kernels using **nvprof** (arch  $\leq 7.x$ ).

**Example:** **matrixprod-simple** using **nsight-sys (nsys-ui)** (cuda  $\geq 11.x$ ).

## Profiling Cuda code

Since Cuda 5 there is a commandline profiler **nvprof**, now being replaced (Cuda 11) with **nsys nvprof**. Use as: **nvprof myprogram <arguments>**

There are also GUI profilers named **nvvp** since Cuda 4 and one named **nsys-ui** since Cuda 11 that replaces it.

**Example:** **matrixprod-simple** with 2 kernels using **nvprof** (arch  $\leq 7.x$ ).

**Example:** **matrixprod-simple** using **nsight-sys** (**nsys-ui**) (cuda  $\geq 11.x$ ).

**Example:** Show **parbody-psm1** (1024grid.par) with the **nsight-sys** GUI.



## Profiling Cuda code

Since Cuda 5 there is a commandline profiler **nvprof**, now being replaced (Cuda 11) with **nsys nvprof**. Use as: **nvprof myprogram <arguments>**

There are also GUI profilers named **nvvp** since Cuda 4 and one named **nsys-ui** since Cuda 11 that replaces it.

**Example:** **matrixprod-simple** with 2 kernels using **nvprof** (arch  $\leq 7.x$ ).

**Example:** **matrixprod-simple** using **nsight-sys** (**nsys-ui**) (cuda  $\geq 11.x$ ).

**Example:** Show **parbody-psm1** (1024grid.par) with the **nsight-sys** GUI.

**Demo:** Show **nvprof** and possibly **nsys-ui** on a modern system (arch  $\leq 7.x$ ).

**Note:** Select **localhost connection**, set commandline with arguments (full path!), enable **Collect CUDA trace** and press **Start**. Select **Timeline View** and expand the **CUDA HW** section to view kernel statistics like in **nvvp**.

Use **zoom** if required (select region, press right mouse and choose **Zoom into selection**).

## Some nice examples

Let's try a few nice (graphic) examples, such as:

- ▶ fluidsGL
- ▶ smokeParticles
- ▶ particles
- ▶ Mandelbrot

## Part 3: Just a simple exercise

## Exercise

**Example:** Let's first shortly discuss the example `MatVecProd-Blas-GPU`, where we compute a matrix vector product using the CuBlas library. So, no need to write your own kernels;-)

## Exercise

**Example:** Let's first shortly discuss the example `MatVecProd-Blas-GPU`, where we compute a matrix vector product using the CuBlas library. So, no need to write your own kernels;-)

**Exercise:** In the directory `MatrixProd-Blas-GPU-exercise` you will find **exactly** the same code as in `MatVecProd-Blas-GPU`, only the names are changed into a matrix matrix product. Your task is to try to modify this example in such a way that it computes a real Matrix Product using CuBlas on the GPU!

**Note:** Maybe the best option is to use the one with explicit memory transfers as there it is much easier to time the actual computation. With unified memory - albeit slightly easier - it is much harder to avoid measuring unexpected implicit and hidden host <-> card transfers as well.

## Exercise

**Example:** Let's first shortly discuss the example `MatVecProd-Blas-GPU`, where we compute a matrix vector product using the CuBlas library. So, no need to write your own kernels;-)

**Exercise:** In the directory `MatrixProd-Blas-GPU-exercise` you will find **exactly** the same code as in `MatVecProd-Blas-GPU`, only the names are changed into a matrix matrix product. Your task is to try to modify this example in such a way that it computes a real Matrix Product using CuBlas on the GPU!

**Note:** Maybe the best option is to use the one with explicit memory transfers as there it is much easier to time the actual computation. With unified memory - albeit slightly easier - it is much harder to avoid measuring unexpected implicit and hidden host <-> card transfers as well.

Discuss shortly the solution.

THE END

Thanks for your attention!